

Literature Review: “Reel or Real? A Context-Aware Deepfake Detection and Verification Framework” pt.3

1. Introduction

Deepfake technology has quickly advanced from early GAN-based models to sophisticated diffusion and transformer architectures that can generate visually and acoustically convincing synthetic media. As generative models become more accessible and high-quality, the focus has shifted from just spotting anomalies to creating reliable verification systems that work well despite real-world distortions, adversarial manipulation, and distribution across different platforms. Recent research in adversarial defences, forensic fingerprinting, perceptual hashing, and ensemble detection shows a clear move from accuracy-focused academic benchmarks to a broader concern with resilience, provenance, and operational use. This literature review explores this development and summarises important technological, forensic, and societal insights to build a foundation for a context-aware, multi-layer deepfake verification framework.

2. Public Awareness & the Social Dimension

Alongside technical progress, recent research highlights the increasing harm caused by deepfakes across personal, organisational, and societal spheres. Studies on the societal effects of deepfakes emphasise the deterioration of trust in digital communication, the spread of political misinformation, and the growth of personalised harassment and reputation damage. Public awareness remains limited: although people express widespread concern about synthetic media, most lack the ability to reliably identify manipulated content, especially when it is distributed through compressed or platform-processed videos. This gap between threat and public ability heightens reliance on verification systems that are transparent, capable of explaining themselves, and scalable. Consequently, the social aspect underscores the importance of robust detectors that consistently perform under real-world conditions and offer understandable signals to non-expert users.

3. Evolution of Deepfake Detection Techniques

Early deepfake detection initially relied heavily on convolutional neural networks trained to recognise spatial artefacts like blending seams, inconsistent lighting, or unnatural textures. These models achieved high accuracy on curated datasets but often performed poorly when faced with unseen manipulation methods, new generator types, or videos affected by compression and noise. Later, transformer-

based architectures and spatio-temporal hybrid models enhanced generalisation by capturing long-range temporal cues, multi-frame inconsistencies, and subtle motion irregularities.

The research landscape then broadened in two main directions. The first includes proactive defences such as adversarial perturbations that endure model retraining and universal adversarial watermarks (e.g., cross-model universal perturbations) designed to diminish deepfake quality at the point of generation. The second direction emphasises robustness through diversity, with ensemble detectors like Deepfake Sentry combining specialised sub-models to address the weaknesses of any individual architecture. These developments signify a shift from point-solution classifiers to multi-layered detection systems intended to withstand both generating variation and adversarial pressure.

4. Dataset Diversity & the Realism Challenge

A recurring observation in recent literature is the decline in detector performance when applied to real-world environments. Benchmark datasets often use studio-quality footage with limited demographic, environmental, and linguistic variation. Consequently, detectors trained on such data frequently perform poorly on low-resolution, surveillance-style footage, social-media compression, or culturally diverse expressions and physiognomies.

Work on CCTV-specific deepfake detection underscores the importance of domain-adapted architectures that can comprehend highly compressed, noisy, and motion-blurred scenes. Likewise, studies on adversarial robustness reveal that even small distribution shifts—such as variations in lighting, camera angles, or compression profiles—can undermine forensic traces that detectors rely on. The literature, therefore, emphasises the need for high-variance, cross-generator, and cross-platform datasets, as well as training regimes that explicitly incorporate real-world distortions rather than assuming pristine data pipelines.

5. Explainability & Trust in AI Detection

As detection architectures become increasingly complex, explainability has become crucial for maintaining public trust, ensuring legal validity, and facilitating deployment at the platform level. Forensic methods such as residual fingerprinting analyse subtle, model-specific artefacts left by generative networks, offering a more understandable basis for decisions than opaque latent features. Perceptual hashing systems—like Hash Shield’s separable hashing architecture—provide compact,

transformation-resistant signatures that help verify media integrity even after resizing, transcoding, or compression.

These methods extend explainability beyond heatmaps and confidence scores by incorporating provenance-aware cues, allowing users and institutions to verify whether the media matches an expected fingerprint or perceptual signature. Such hybrid verification mechanisms bridge technical detection with human interpretability, enabling workflows where detection outcomes can be audited, validated, and legally scrutinised. The literature consistently highlights that trust is not built solely on accuracy, but on transparency, reproducibility, and provenance.

6. Emerging Approaches & Gaps

Across the reviewed literature, various promising approaches arise:

- Proactively disrupt deepfake creation with training-resistant adversarial perturbations and universal watermarks that prevent it before deployment.
- Forensic fingerprinting and perceptual hashing, offering stable signatures for tamper detection and provenance verification.
- Ensemble and domain-specific detectors, which enhance robustness across environments through architectural diversity and context-aware training.
- Research on adversarial attacks exposing how existing detectors are vulnerable to transferable, imperceptible perturbations created via gradient decomposition or surrogate-model strategies.

Despite these advancements, significant gaps still exist. Most techniques are tested in isolation rather than as part of comprehensive verification systems. Multimodal robustness (e.g., combining video, audio, metadata, and hashing signals) remains insufficiently explored. Adversarial robustness is often evaluated only after the model is deployed, rather than being incorporated into the detection architecture from the outset. Lastly, usability and explainability for real-world stakeholders—such as journalists, educators, legal experts, and everyday users—receive limited attention. These gaps highlight the necessity for systems that integrate prevention, detection, provenance, and interpretability within unified, adversarially aware frameworks.

7. Conclusion & Research Justification

The current literature shows a maturing yet fragmented field shifting from classifier-focused detection to multi-layered verification. Initiative-taking watermarking, residual fingerprints, perceptual hashing, ensemble pipelines, and adversarial stress-testing each add valuable elements, but no existing approach combines

them into a coherent system capable of withstanding real-world distribution shifts and adaptive attacks.

This review highlights the necessity of a context-aware deepfake video verification framework that combines proactive defences, forensic integrity measures, and robust multi-modal detection within a transparent, user-focused design. Such a system would not only identify synthetic media but also provide verifiable provenance signals, resilience against attacks, and interpretable outputs suitable for high-stakes environments. By addressing the limitations identified in existing research—such as lack of integration, weak adversarial design, limited multimodality, and insufficient focus on explainability—this project aims to develop a comprehensive, deployable solution for digital authenticity in the age of synthetic media.

Building on these insights, this project proposes a software-engineering-oriented, context-aware deepfake detection and verification framework that:

- Combines multimodal data and cryptographic provenance to ensure end-to-end authenticity verification.
- Offers explainable outputs like heatmaps, confidence scores, and provenance chains to build user trust.
- Emphasises cross-dataset generalisation via hybrid model design and real-world testing.
- Incorporates ethical, user-focused principles into its design for accessibility and transparency.

By integrating machine learning, human-centred design, and digital ethics, this work advances both technically and socially to restore authenticity in the age of synthetic media, aligning with the academic and moral imperatives that define 21st-century AI research.

8. References

Cui, J., Yu, B., Wang, Q., Meng, F. and Lu, J. 2025. WhiADD: Semantic-Acoustic Fusion for Robust Audio Deepfake Detection, Proceedings of the 33rd ACM International Conference on Multimedia (MM '25), 27–31 October 2025, pp. 1–9.

Jiang, J., Li, B., Yu, S., Liu, C., An, S., Liu, M. and Yu, M. 2021. A Residual Fingerprint-Based Defense Against Adversarial Deepfakes, 2021 IEEE 23rd International Conference on High Performance Computing & Communications (HPCC) et al. (co-located conferences), pp. 797–804.

Yang, M., Qi, B., Ma, R., Xian, Y. and Ma, B. 2025. HashShield: A Robust DeepFake Forensic Framework With Separable Perceptual Hashing, *IEEE Signal Processing Letters*, vol. 32, no. 12, pp. 1186–1190.

Göller, N., Graner, L., Frick, R.A. and Bunzel, N. 2025. Team RoMa @ AADD-2025: On the Generation of Transferable and Visually Imperceptible Adversarial Attacks Against Deepfake Detectors, *Proceedings of the 33rd ACM International Conference on Multimedia (MM '25)*, 27–31 October 2025, pp. 1–5.

Segalis, E. and Galili, E. 2020. OGAN: Disrupting Deepfakes with an Adversarial Attack that Survives Training, *arXiv preprint arXiv:2006.12247*.

Seo, W., Baek, J., Jung, Y. and Park, S. 2025. MIG-COW: Transferable Adversarial Attacks on Deepfake Detectors via Gradient Decomposition, *Proceedings of the 33rd ACM International Conference on Multimedia (MM '25)*, 27–31 October 2025, pp. 1–7.

Ștefan, L.-D., Stanciu, D.-C., Dogariu, M., Constantin, M.G., Jitaru, A.C. and Ionescu, B. 2023. Deepfake Sentry: Harnessing Ensemble Intelligence for Resilient Detection and Generalisation, *U.P.B. Scientific Bulletin, Series C (Electrical Engineering and Computer Science)*, vol. 85, no. 4, pp. 201–216.

Imran, J., Bawany, N.Z. and Tahzeeb, S. 2025. Deepfake Detection in CCTV Videos Using Deep Learning, *Arabian Journal for Science and Engineering (Online First, 19 August 2025)*, pp. 1–23.

Huang, H., Wang, Y., Chen, Z., Zhang, Y., Li, Y., Tang, Z., Chu, W., Chen, J., Lin, W. and Ma, K.-K. 2022. CMUA-Watermark: A Cross-Model Universal Adversarial Watermark for Combating Deepfakes, *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 36, no. 1, pp. 989–997.

Le, T.-N., Nguyen, H.H., Yamagishi, J. and Echizen, I. 2022. Robust Deepfake on Unrestricted Media: Generation and Detection, in M. Khosravy et al. (eds) *Frontiers in Fake Media Generation and Detection, Studies in Autonomic, Data-driven and Industrial Computing*, Springer, pp. 81–108.